

# Comprehensive Embedded AI & Signal Processing Notes

## Table of Contents

1. Machine Learning Hardware (MAC, ALU, Accumulator)
  2. ARM Cortex-M55 Processor
  3. Ethos-U MicroNPUs
  4. CMSIS-NN
  5. ADC and Quantization
  6. Sampling Theory (Nyquist Criteria)
  7. Signal Processing (FFT, Fourier Transform)
  8. Filters (Notch Filter)
  9. Audio Processing (MFCC, Mel Frequency)
  10. ATmega328P Microcontroller
  11. Confusion Matrix Calculations
  12. Embedded AI Deployment
- 

## 1. Machine Learning Hardware

### 1.1 ALU vs MAC vs Accumulator

#### ALU (Arithmetic Logic Unit)

- **Function:** Performs basic arithmetic (addition, subtraction) and logical operations (AND, OR, XOR)
- **Operations:** Single-cycle operations
- **Use Case:** General-purpose computing, control flow
- **Example:**

$$A = 5 + 3 = 8$$

$$B = 10 - 2 = 8$$

$$C = A \text{ AND } B = 8 \text{ AND } 8 = 8$$

#### MAC (Multiply-Accumulate Unit)

- **Function:** Performs multiplication followed by addition in a single operation
- **Operations:**  $\text{Result} = (A \times B) + C$
- **Use Case:** Neural networks, DSP, matrix operations
- **Efficiency:** Critical for ML - most NN operations are dot products

### **Example Calculation:**

Neural Network Weight Calculation:

Weights: [0.5, 0.3, 0.2]

Inputs: [1.0, 2.0, 3.0]

MAC Operation:

$$\text{Result} = (0.5 \times 1.0) + (0.3 \times 2.0) + (0.2 \times 3.0)$$

$$= 0.5 + 0.6 + 0.6$$

$$= 1.7$$

Traditional ALU would need:

$$\text{Step 1: } 0.5 \times 1.0 = 0.5$$

$$\text{Step 2: } 0.3 \times 2.0 = 0.6$$

$$\text{Step 3: } 0.2 \times 3.0 = 0.6$$

$$\text{Step 4: } 0.5 + 0.6 = 1.1$$

$$\text{Step 5: } 1.1 + 0.6 = 1.7$$

(5 operations vs 1 MAC operation)

### **Accumulator**

- **Function:** Special register that stores intermediate results
- **Role:** Holds the sum in MAC operations
- **Width:** Determines precision (8-bit, 16-bit, 32-bit, etc.)

## **1.2 Accumulator Width Impact**

<b>Accumulator Width</b>	<b>Range</b>	<b>ML Impact</b>
<b>8-bit</b>	0 to 255 (unsigned)	Limited precision, overflow risk, suitable for simple models
<b>16-bit</b>	0 to 65,535	Better for quantized networks
<b>32-bit</b>	0 to 4,294,967,295	High precision, prevents overflow in deep networks

### **Example:**

16-bit Accumulator:

Max value = 65,535

If we have 1000 MAC operations:

Average value per operation =  $65,535 / 1000 = 65.5$

This limits weight  $\times$  input product to ~65

32-bit Accumulator:

Max value = 4,294,967,295

Average per 1000 ops = 4,294,967

Much more headroom for complex calculations

### 1.3 Ethos-U55 MAC Engine Configurations

Configuration	MAC Units	Performance	Use Case
<b>32 MAC</b>	32	Lowest power	Keyword spotting, simple inference
<b>64 MAC</b>	64	Balanced	Face detection, gesture recognition
<b>128 MAC</b>	128	High performance	Object detection
<b>256 MAC</b>	256	Maximum performance	Real-time video processing

## 2. ARM Cortex-M55 Processor

### 2.1 Key Features

#### Cortex-M55 Specifications:

- **First CPU with ARM Helium Technology** (M-profile Vector Extension - MVE)
- **Performance Gains:**
  - Up to **15x ML performance** vs previous Cortex-M generations
  - Up to **5x DSP performance**
- **Vector Processing:** SIMD (Single Instruction Multiple Data)
- **TrustZone Support:** Hardware-level security
- **Power Efficiency:** Designed for battery-powered devices

### 2.2 Helium Technology

#### What is Helium?

- Vector processing extension for Cortex-M
- Processes multiple data elements in parallel
- Optimized for ML and DSP workloads

#### Example:

Traditional Processing (without Helium):

Add 4 numbers: [1, 2, 3, 4] + [5, 6, 7, 8]

Cycle 1: 1 + 5 = 6

Cycle 2:  $2 + 6 = 8$

Cycle 3:  $3 + 7 = 10$

Cycle 4:  $4 + 8 = 12$

Total: 4 cycles

With Helium (SIMD):

Cycle 1:  $[1,2,3,4] + [5,6,7,8] = [6,8,10,12]$

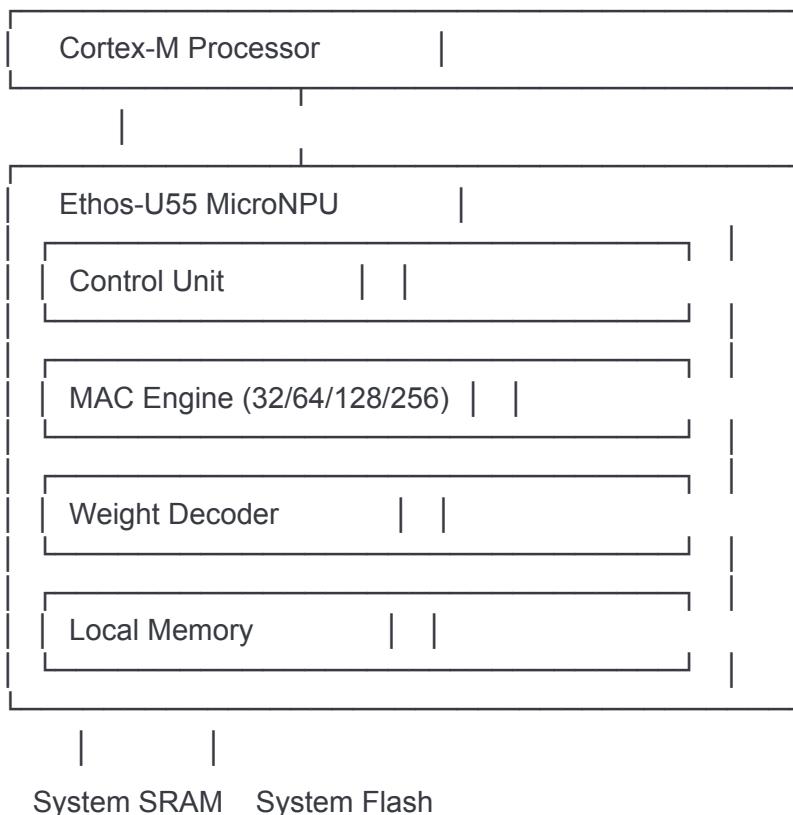
Total: 1 cycle (4x faster!)

---

### 3. Ethos-U MicroNPUs

#### 3.1 Ethos-U55 (First Generation)

**Architecture:**



**Key Features:**

- **MAC Options:** 32, 64, 128, or 256 units
- **Weight Compression:** On-the-fly decompression (up to 90% SRAM reduction)
- **Compatible Processors:** Cortex-M55, M33, M7, M4
- **Memory:** Works with SRAM and Flash

## 3.2 Ethos-U65 (Second Generation)

### Improvements over U55:

- **MAC Options:** 256 or 512 units ( $2\times$  more powerful)
- **DRAM Support:** Better for larger models
- **Hybrid Systems:** Can work in A-class systems (Cortex-A processors)

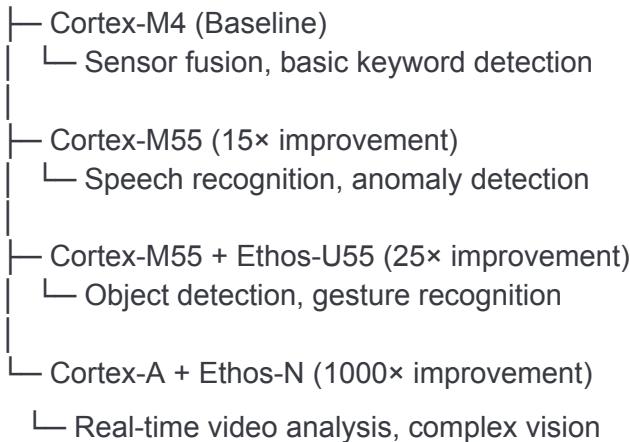
## 3.3 Suitable Applications

### Ethos-U55 Applications (From Exam Q2)

1. **BLE Sense Nano (Cortex-M55 + Ethos-U55)**
  - **Object Classification:** Identify objects in images
  - **Why suitable?**
    - Low power for battery operation
    - On-device processing (privacy)
    - Fast inference (<100ms)
2. **Real-world Examples:**
  - **Keyword Spotting:** "Hey Google", "Alexa"
  - **Predictive Maintenance:** Vibration analysis in motors
  - **Gesture Recognition:** Smart home controls
  - **Face Unlock:** Smartphone security

### Application Mapping by Performance

Performance Scale:



---

## 4. CMSIS-NN

### 4.1 What is CMSIS-NN?

#### CMSIS (Cortex Microcontroller Software Interface Standard)

- Software library for Cortex-M processors

- **CMSIS-NN:** Neural Network optimized kernels
- Provides efficient implementations of common NN operations

## 4.2 Performance Comparison

**From Document - Wav2Letter Performance:**

Baseline (CM4 Reference):	1×
CM4 with CMSIS:	3.5×
Cortex-M55 with CMSIS:	11×
Cortex-M55 + Ethos-U55:	25×

Total Improvement: ~1000× from baseline to U55!

## 4.3 TensorFlow Lite Micro Integration

**Workflow:**

1. Train Model (TensorFlow)
  2. Convert to .tflite file
  3. Optimize with Vela Compiler
  4. Deploy to device with TFLite Micro
  5. Run inference using:
    - Reference kernels (slow)
    - CMSIS-NN kernels (fast)
    - Ethos-U microNPU (fastest)
- 

# 5. ADC and Quantization

## 5.1 N-bit ADC Resolution

**Formula:**

Number of Quantization Levels =  $2^N$

Voltage Resolution =  $(V_{max} - V_{min}) / (2^N)$

## 5.2 10-bit ADC Example (Step-by-Step)

**Given:**

- Reference Voltage: 5V
- ADC: 10-bit
- Input Signal: 3.2V

### **Step 1: Calculate Quantization Levels**

Levels =  $2^{10} = 1024$  levels

### **Step 2: Calculate Resolution (LSB)**

$$\begin{aligned}\text{Resolution} &= V_{\text{ref}} / \text{Levels} \\ &= 5V / 1024 \\ &= 0.00488V \text{ (4.88mV per step)}\end{aligned}$$

### **Step 3: Convert Analog to Digital**

$$\begin{aligned}\text{Digital Value} &= (V_{\text{input}} / V_{\text{ref}}) \times (2^N - 1) \\ &= (3.2V / 5V) \times 1023 \\ &= 0.64 \times 1023 \\ &= 654.72 \approx 655 \text{ (rounded)}\end{aligned}$$

Binary: 655 = 1010001111 (10 bits)

### **Step 4: Quantization Error**

$$\text{Actual Voltage} = (655 / 1023) \times 5V = 3.201V$$

$$\text{Quantization Error} = 3.201V - 3.2V = 0.001V = 1mV$$

## **5.3 Different ADC Resolutions**

ADC Bits	Level s	Resolution (5V)	Accuracy
8-bit	256	19.5mV	$\pm 9.75mV$
10-bit	1024	4.88mV	$\pm 2.44mV$
12-bit	4096	1.22mV	$\pm 0.61mV$
16-bit	65,536	76.3 $\mu$ V	$\pm 38.1\mu$ V

## **5.4 Linear Quantization**

### **Process:**

1. Divide voltage range into equal intervals
2. Map each interval to a digital code
3. Round to nearest level

### **Example: 3-bit ADC (8 levels), 0-8V range**

Level 0: 0.0 - 1.0V → 000

Level 1: 1.0 - 2.0V → 001  
Level 2: 2.0 - 3.0V → 010  
Level 3: 3.0 - 4.0V → 011  
Level 4: 4.0 - 5.0V → 100  
Level 5: 5.0 - 6.0V → 101  
Level 6: 6.0 - 7.0V → 110  
Level 7: 7.0 - 8.0V → 111

Input: 5.7V → Level 5 → Binary: 101

## 5.5 Brain Signal Interfacing (EEG)

### From Exam Q1: EEG Signal Processing

**Given:**

- EEG Signal: Up to ~70 Hz
- Application: Brain-Computer Interface

### Step 1: Apply Nyquist Criterion

$$\begin{aligned}\text{Minimum Sampling Rate} &= 2 \times f_{\max} \\ &= 2 \times 70 \text{ Hz} \\ &= 140 \text{ Hz}\end{aligned}$$

Practical Sampling Rate: 256 Hz or 512 Hz

(Power of 2 for FFT efficiency)

### Step 2: Choose ADC Resolution

EEG Signal Amplitude: ~10µV to 100µV (typical)  
Noise Level: ~1µV

Required SNR: 40-60 dB

For 12-bit ADC:

$$\text{Dynamic Range} = 20 \times \log_{10}(2^{12}) = 72 \text{ dB } \checkmark$$

Recommended: 12-bit or 16-bit ADC

### Step 3: Calculate Data Rate (16 channels)

$$\begin{aligned}\text{Data Rate} &= \text{Sampling Rate} \times \text{Bits} \times \text{Channels} \\ &= 256 \text{ Hz} \times 16 \text{ bits} \times 16 \text{ channels} \\ &= 65,536 \text{ bits/second} \\ &= 8.192 \text{ KB/second}\end{aligned}$$

## 5.6 Non-Linear Quantization

### Why Non-Linear?

- Human perception is logarithmic (hearing, vision)
- More resolution for small signals, less for large

### $\mu$ -law Encoding (Audio):

Formula:  $F(x) = \text{sign}(x) \times \ln(1 + \mu|x|) / \ln(1 + \mu)$

Where  $\mu = 255$  (North America) or  $256$  (International)

Example:

Input:  $0.1$  (normalized)

Output:  $\text{sign}(0.1) \times \ln(1 + 255 \times 0.1) / \ln(256)$

$$= \ln(26.5) / \ln(256)$$

$$= 0.587$$

This gives more bits to quiet sounds!

### A-law Encoding (Used in Europe):

- Similar to  $\mu$ -law but different curve
  - Better for small amplitude signals
- 

## 6. Sampling Theory

### 6.1 Nyquist Criterion

Theorem:

$$f_{\text{sampling}} \geq 2 \times f_{\text{max}}$$

Where  $f_{\text{max}}$  is the highest frequency in the signal

### Why 2×?

- Need at least 2 samples per cycle to reconstruct signal
- Prevents aliasing (frequency folding)

### 6.2 Practical Example

Audio Signal: 20 Hz - 20 kHz

Minimum Sampling Rate =  $2 \times 20,000$  Hz = 40 kHz

Actual CD Quality = 44.1 kHz (10% overhead)

### **EEG Signal: 0.5 Hz - 70 Hz**

Minimum =  $2 \times 70$  Hz = 140 Hz

Practical = 256 Hz (allows filtering)

## **6.3 Aliasing Effect**

### **The "Wagon Wheel Effect"**

- Wheel spins forward but appears to move backward
- Happens when sampling rate <  $2 \times$  rotation frequency

#### **Example:**

Wheel rotating at 10 Hz (10 revolutions/second)

Camera at 15 FPS

$15 \text{ FPS} < 2 \times 10 \text{ Hz} = 20 \text{ Hz}$  (violates Nyquist!)

Result: Appears to rotate backward

Solution: Film at  $\geq 20$  FPS

### **Mathematical Explanation:**

True frequency:  $f = 10$  Hz

Sampling:  $fs = 15$  Hz

Aliased frequency =  $|f - fs| = |10 - 15| = 5$  Hz

Appears as 5 Hz backward rotation!

---

## **7. Signal Processing**

### **7.1 Fourier Transform**

**Purpose:** Convert time-domain signal to frequency-domain

#### **Time Domain → Frequency Domain**

Signal:  $x(t) = A \times \sin(2\pi ft)$

Fourier Transform gives:

- Frequency components
- Amplitude of each frequency
- Phase of each frequency

#### **Example:**

Mixed Signal:  $x(t) = 2 \times \sin(2\pi \times 50t) + 1 \times \sin(2\pi \times 120t)$

Time Domain: Complex waveform

Frequency Domain:

- Peak at 50 Hz (amplitude = 2)
- Peak at 120 Hz (amplitude = 1)

## **7.2 FFT (Fast Fourier Transform)**

#### **Efficiency:**

- DFT:  $O(N^2)$  operations
- FFT:  $O(N \log N)$  operations

#### **For N = 1024 samples:**

DFT:  $1,024^2 = 1,048,576$  operations

FFT:  $1024 \times \log_2(1024) = 1024 \times 10 = 10,240$  operations

FFT is 102× faster!

#### **Practical Implementation:**

Input: 512 samples at 256 Hz

FFT Size: 512 (must be power of 2)

Frequency Resolution = Sampling Rate / FFT Size

$$= 256 \text{ Hz} / 512$$

$$= 0.5 \text{ Hz per bin}$$

Output: 256 frequency bins (0-128 Hz)

## **7.3 Determining Signal Frequencies**

#### **Step-by-Step Process:**

#### **Given Signal:**

$$x(t) = 3 \times \sin(2\pi \times 10t) + 2 \times \sin(2\pi \times 25t) + 0.5 \times \sin(2\pi \times 50t)$$

### **Step 1: Identify Components**

- Component 1:  $f_1 = 10 \text{ Hz}, A_1 = 3$
- Component 2:  $f_2 = 25 \text{ Hz}, A_2 = 2$
- Component 3:  $f_3 = 50 \text{ Hz}, A_3 = 0.5$

### **Step 2: Determine $f_{\max}$**

$f_{\max} = \text{maximum frequency} = 50 \text{ Hz}$

### **Step 3: Required Sampling Rate**

$f_s = 2 \times f_{\max} = 2 \times 50 = 100 \text{ Hz}$  minimum

Practical: 128 Hz or 256 Hz

---

## **8. Filters**

### **8.1 Notch Filter**

**Purpose:** Remove specific frequency (noise)

**Common Applications:**

- Remove 50/60 Hz power line noise
- Remove specific interference

**Design:**

Center Frequency ( $f_0$ ): Frequency to remove

Bandwidth (BW): Range around  $f_0$

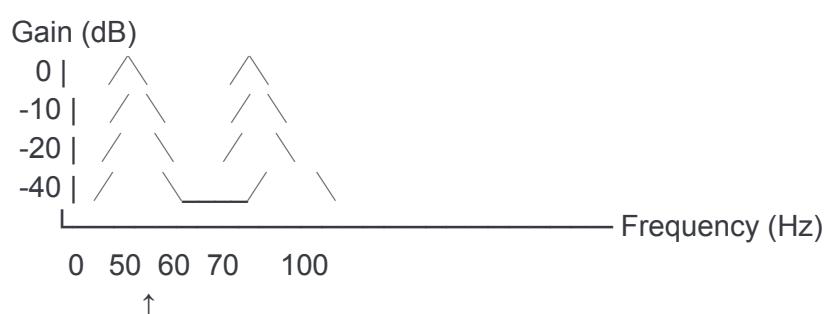
Quality Factor (Q):  $f_0 / \text{BW}$

Example:

Remove 60 Hz with 2 Hz bandwidth

$Q = 60 / 2 = 30$  (narrow notch)

**Frequency Response:**



Notch at 60 Hz

## 8.2 Other Common Filters

### Low-Pass Filter:

- Passes frequencies below cutoff
- Use: Remove high-frequency noise

### High-Pass Filter:

- Passes frequencies above cutoff
- Use: Remove DC offset, low-frequency drift

### Band-Pass Filter:

- Passes frequencies in a range
  - Use: Isolate specific frequency band
- 

## 9. Audio Processing

### 9.1 Mel Frequency Scale

#### Why Mel Scale?

- Human hearing is logarithmic, not linear
- We perceive pitch differences logarithmically

#### Conversion:

$$\text{Mel}(f) = 2595 \times \log_{10}(1 + f/700)$$

Examples:

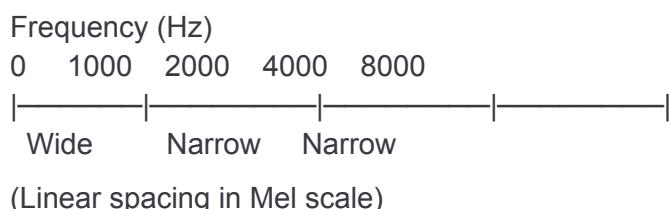
$$f = 0 \text{ Hz} \rightarrow \text{Mel} = 0$$

$$f = 1000 \text{ Hz} \rightarrow \text{Mel} = 1127$$

$$f = 2000 \text{ Hz} \rightarrow \text{Mel} = 1842$$

$$f = 4000 \text{ Hz} \rightarrow \text{Mel} = 2555$$

#### Mel Filter Bank:



## 9.2 MFCC (Mel-Frequency Cepstral Coefficients)

**Purpose:** Feature extraction for speech recognition

**Process:**

1. Frame signal (20-40ms windows)
2. Apply FFT → Frequency spectrum
3. Apply Mel filter bank → Mel spectrum
4. Take logarithm → Log Mel spectrum
5. Apply DCT → MFCC features

**Step-by-Step Example:**

**Given:** Speech signal at 16 kHz

Step 1: Frame

Window size = 25ms = 0.025s

Samples per frame =  $16,000 \times 0.025 = 400$  samples

Step 2: FFT

FFT size = 512 (next power of 2)

Output: 256 frequency bins (0-8 kHz)

Step 3: Mel Filter Bank (40 filters)

Output: 40 Mel-scaled values

Step 4: Log

Output: 40 log-scaled values

Step 5: DCT

Keep first 13 coefficients (MFCC 1-13)

Final Output: 13 features per frame

**Why MFCCs are Powerful:**

- Compact representation (13 numbers instead of 400)
- Captures timbral characteristics of speech
- Robust to noise

---

## 10. ATmega328P Microcontroller

### 10.1 Overview

## Specifications:

- **Architecture:** 8-bit AVR RISC
- **Clock:** Up to 20 MHz
- **Flash:** 32 KB
- **RAM:** 2 KB
- **EEPROM:** 1 KB
- **Package:** DIP-28, TQFP-32, QFN-32

## 10.2 Pin Configuration (DIP-28 Package)

PC6	1	RESET	VCC	28	VCC (5V)
PD0	2	RXD	GND	27	GND
PD1	3	TXD	XTAL	26	PC5 (ADC5/SCL)
PD2	4		XTAL	25	PC4 (ADC4/SDA)
PD3	5		PC3	24	PC3 (ADC3)
PD4	6		PC2	23	PC2 (ADC2)
VCC	7		PC1	22	PC1 (ADC1)
GND	8		PC0	21	PC0 (ADC0)
PB6	9		GND	20	GND
PB7	10		AREF	19	AREF
PD5	11		PB5	18	PB5 (SCK)
PD6	12		PB4	17	PB4 (MISO)
PD7	13		PB3	16	PB3 (MOSI)
PB0	14		PB2	15	PB2

## 10.3 Pin Functions

### Port B (PB0-PB7)

- **PB0-PB5:** Digital I/O
- **PB3-PB5:** SPI (MOSI, MISO, SCK)
- **PB6-PB7:** Crystal oscillator

### Port C (PC0-PC5)

- **PC0-PC5:** Digital I/O + ADC (6 channels)
- **PC4-PC5:** I2C (SDA, SCL)
- **PC6:** RESET pin

### Port D (PD0-PD7)

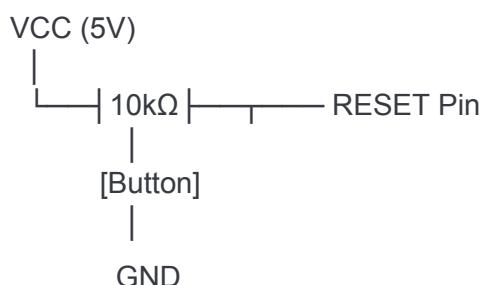
- **PD0-PD1:** UART (RX, TX)
- **PD2-PD3:** External interrupts (INT0, INT1)
- **PD3, PD5, PD6:** PWM outputs
- **PD4, PD7:** Timer inputs

## 10.4 RESET Pin Function

### PC6/RESET Pin:

- **Active Low:** Pulling to ground resets MCU
- **Internal Pull-up:** Normally HIGH
- **Reset Sources:**
  1. Power-on Reset
  2. External Reset (button)
  3. Watchdog Reset
  4. Brown-out Reset

### Reset Circuit:



## 10.5 Clock System

### Internal Oscillator:

- 8 MHz factory calibrated
- Can be prescaled to lower frequencies

### External Crystal:

- Up to 20 MHz
- More accurate for timing-critical applications

### Clock Frequency Calculations:

Timer Frequency = Clock / Prescaler

Example: 16 MHz clock with prescaler 64

Timer Freq =  $16,000,000 / 64 = 250,000 \text{ Hz}$

Time per tick =  $1 / 250,000 = 4 \mu\text{s}$

## 10.6 Quad Package (QFN/TQFP)

### Advantages:

- Smaller footprint

- Better thermal performance
- More pins (32 vs 28)

#### **Additional Pins in 32-pin package:**

- More power/ground pins
  - Additional Port C pins
  - Better noise immunity
- 

## **11. Confusion Matrix Calculations**

### **11.1 Understanding the Matrix**

#### **From Exam Q3:**

Confusion Matrix (Motion Detection):

		Predicted				
		Circle	Idle	Left-Right	Up-Down	
Actual	Circle	205	10	1	46	
	Idle	6	199	0	32	
Left-Right		9	17	223	34	
Up-Down		21	8	3	186	

### **11.2 Calculate Metrics for Each Class**

#### **Circle Class:**

**True Positive (TP):** 205 (correctly classified as Circle)

#### **False Positive (FP):**

$$\begin{aligned}
 \text{FP} &= \text{Predicted Circle but actually other classes} \\
 &= 6 (\text{Idle}) + 9 (\text{Left-Right}) + 21 (\text{Up-Down}) \\
 &= 36
 \end{aligned}$$

#### **False Negative (FN):**

$$\begin{aligned}
 \text{FN} &= \text{Actually Circle but predicted as other} \\
 &= 10 (\text{Idle}) + 1 (\text{Left-Right}) + 46 (\text{Up-Down}) \\
 &= 57
 \end{aligned}$$

#### **True Negative (TN):**

$$\begin{aligned}
 \text{TN} &= \text{All other correct classifications} \\
 &= 199 + 0 + 32 + 223 + 34 + 3 + 186
 \end{aligned}$$

= 677

#### Per-Class Accuracy:

$$\begin{aligned}\text{Accuracy} &= (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN}) \\ &= (205 + 677) / (205 + 677 + 36 + 57) \\ &= 882 / 975 \\ &= 0.9046 = 90.46\%\end{aligned}$$

#### TPR (Recall/Sensitivity):

$$\begin{aligned}\text{TPR} &= \text{TP} / (\text{TP} + \text{FN}) \\ &= 205 / (205 + 57) \\ &= 205 / 262 \\ &= 0.7824 = 78.24\%\end{aligned}$$

#### TNR (Specificity):

$$\begin{aligned}\text{TNR} &= \text{TN} / (\text{TN} + \text{FP}) \\ &= 677 / (677 + 36) \\ &= 677 / 713 \\ &= 0.9495 = 94.95\%\end{aligned}$$

#### PPV (Precision):

$$\begin{aligned}\text{PPV} &= \text{TP} / (\text{TP} + \text{FP}) \\ &= 205 / (205 + 36) \\ &= 205 / 241 \\ &= 0.8506 = 85.06\%\end{aligned}$$

#### F1 Score:

$$\begin{aligned}\text{F1} &= 2 \times (\text{PPV} \times \text{TPR}) / (\text{PPV} + \text{TPR}) \\ &= 2 \times (0.8506 \times 0.7824) / (0.8506 + 0.7824) \\ &= 2 \times 0.6656 / 1.633 \\ &= 0.8151 = 81.51\%\end{aligned}$$

#### Complete Results for All Classes:

Class	TP	FP	FN	TN	Accuracy	TPR	TNR	PPV	F1
Circle	205	36	57	677	90.46%	78.24%	94.95%	85.06%	81.51%
Idle	199	35	38	703	92.51%	83.97%	95.26%	85.04%	84.50%
Left-Right	223	4	60	688	93.44%	78.80%	99.42%	98.24%	87.40%

Up-Down	186	112	32	645	85.23%	85.32%	85.20%	62.42%	72.09%
---------	-----	-----	----	-----	--------	--------	--------	--------	--------

## 11.3 Overall Accuracy

Total TP = 205 + 199 + 223 + 186 = 813

Total Samples = 262 + 237 + 283 + 218 = 1000

Overall Accuracy = Total TP / Total Samples

= 813 / 1000

= 0.813 = 81.3%

---

# 12. Embedded AI Deployment

## 12.1 Deployment Pipeline (From Exam Q4)

### Step 1: Data Acquisition

- Collect representative data
- Label accurately
- Balance classes
- Split: Train (70%), Validation (15%), Test (15%)

### Step 2: Model Training

- Choose architecture (CNN, RNN, etc.)
- Train on desktop/cloud
- Optimize hyperparameters
- Achieve target accuracy

### Step 3: Model Optimization

- Quantization (32-bit → 8-bit)
- Pruning (remove unnecessary weights)
- Knowledge distillation
- Use Vela optimizer (for Ethos-U)

### Step 4: Conversion

- Convert to TensorFlow Lite (.tflite)
- Optimize for embedded target
- Verify accuracy maintained

### Step 5: Deployment

- Flash model to device
- Integrate with application code
- Test on real hardware
- Monitor performance

## **Step 6: Handle Issues**

### **a) Model Fails for Test Data:**

Diagnosis:

- Overfitting to training data
- Insufficient diverse training samples
- Model too complex for task

Solutions:

1. Collect more diverse test data
2. Apply data augmentation
3. Use regularization (dropout, L2)
4. Simplify model architecture
5. Increase validation set

### **b) Low Accuracy on Smartphone + Nano Combined:**

Diagnosis:

- Quantization errors
- Hardware incompatibility
- Different sensor characteristics

Solutions:

1. Re-quantize with representative data from both devices
2. Use quantization-aware training
3. Calibrate sensors
4. Create device-specific models
5. Use transfer learning

## **12.2 Performance Monitoring**

### **Key Metrics:**

1. Inference Time: <100ms for real-time
  2. Power Consumption: <10mW for always-on
  3. Memory Usage: <512KB RAM typical
  4. Accuracy: >90% for production
  5. False Positive Rate: <5%
-

# **Summary: Key Formulas**

## **ADC & Sampling**

Quantization Levels =  $2^N$

Resolution =  $V_{ref} / 2^N$

Sampling Rate  $\geq 2 \times f_{max}$  (Nyquist)

Data Rate = Sampling Rate  $\times$  Bits  $\times$  Channels

## **Confusion Matrix**

Accuracy =  $(TP + TN) / Total$

TPR (Recall) =  $TP / (TP + FN)$

TNR (Specificity) =  $TN / (TN + FP)$

PPV (Precision) =  $TP / (TP + FP)$

F1 Score =  $2 \times (PPV \times TPR) / (PPV + TPR)$

## **Signal Processing**

FFT Resolution = Sampling Rate / FFT Size

Mel(f) =  $2595 \times \log_{10}(1 + f/700)$

Timer Frequency = Clock / Prescaler

---