What is tinyML and what is it used for?

On-device machine learning applications in the single mW and below

Vibration and motion

Any ‘signal’

Predictive maintenance, sensor fusion, accelerometer, pressure, lidar/radar, speed, shock, vibration, pollution, density, viscosity, etc.

Voice and sound

Recognition and creation

Keyword spotting, speech recognition, natural language processing, speech synthesis, sound recognition, etc.

Vision

Images and video

Object detection, face unlock, object classification etc.

arm

Pushing the Boundaries for Real-time On-device Processing

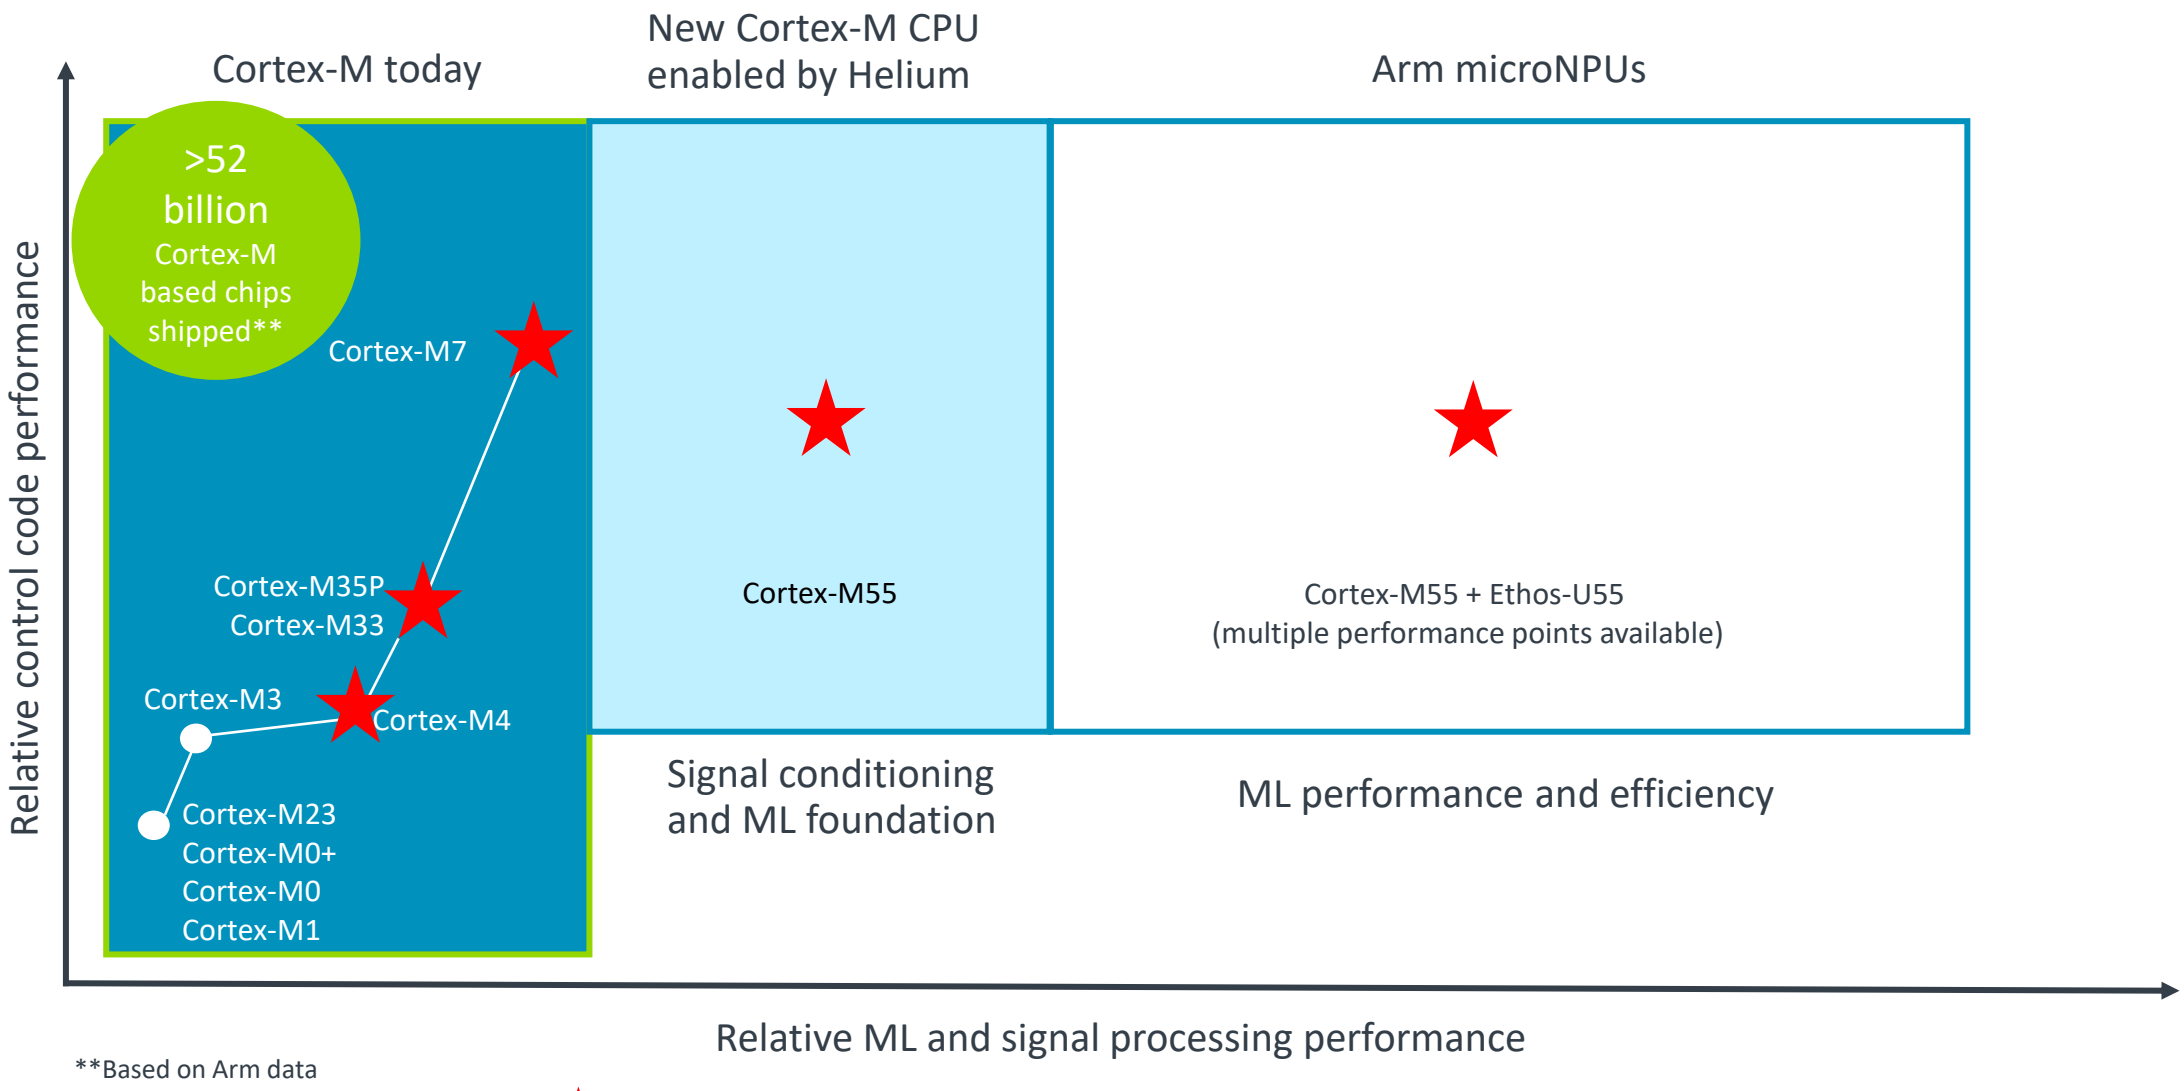

**Based on Arm data

Well suited for ML & DSP applications

Graph not to scale

Cortex-M55: The Most AI-capable Cortex-M Processor

- ✓ First CPU based on Arm Helium technology
  - Energy-efficient and configurable with vector processing capabilities
  - Delivers up to 5x DSP performance and up to 15x ML performance*
  - Versatile capability for both classical ML and NN inference
- ✓ Advanced memory interfaces for fast access to ML data and weights
- ✓ TrustZone support
- ✓ Extensive configurability

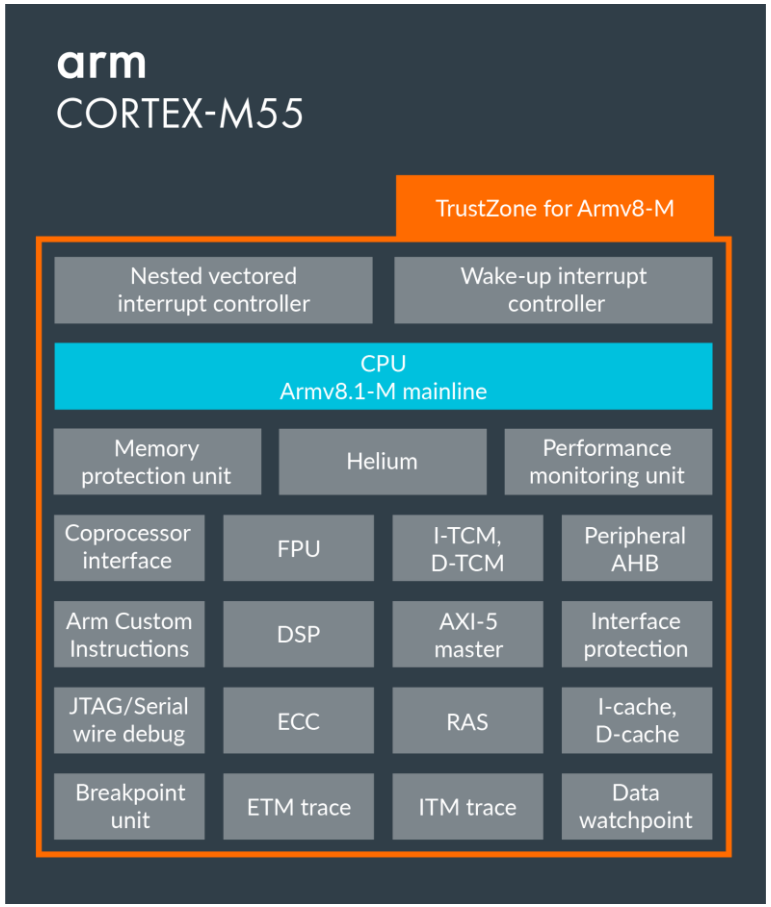

*Compared to previous Cortex-M generations

arm

Ethos-U55: The first microNPU for Cortex-M

- ✓ Highest efficiency and small memory footprint
- ✓ 32, 64, 128, or 256 unit multiply-accumulate (MAC) engine
- ✓ Weight decoder and DMA for on-the-fly weight decompression
- ✓ Tooling available for offline optimization
- ✓ Works with a range of Cortex-M processors:
  - Cortex-M55
  - Cortex-M7
  - Cortex-M33
  - Cortex-M4

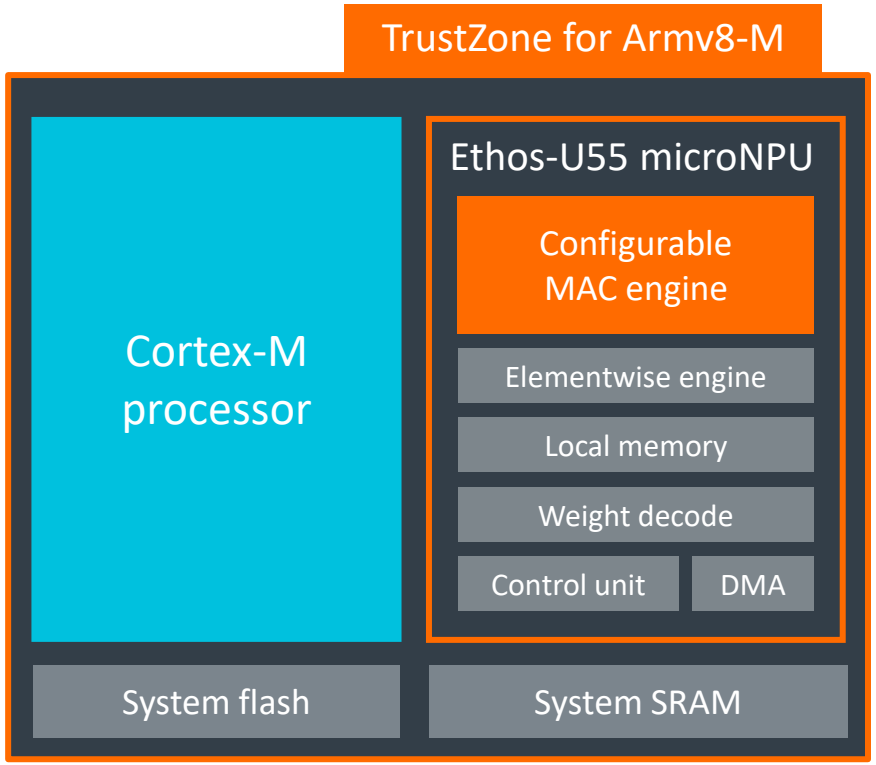

arm

Ethos-U65: The second generation microNPU

- ✓ 256 or 512 unit multiply-accumulate (MAC) engine
- ✓ DMA update for DRAM as well as flash support
- ✓ Can be an M-class subsystem inside an A-class system

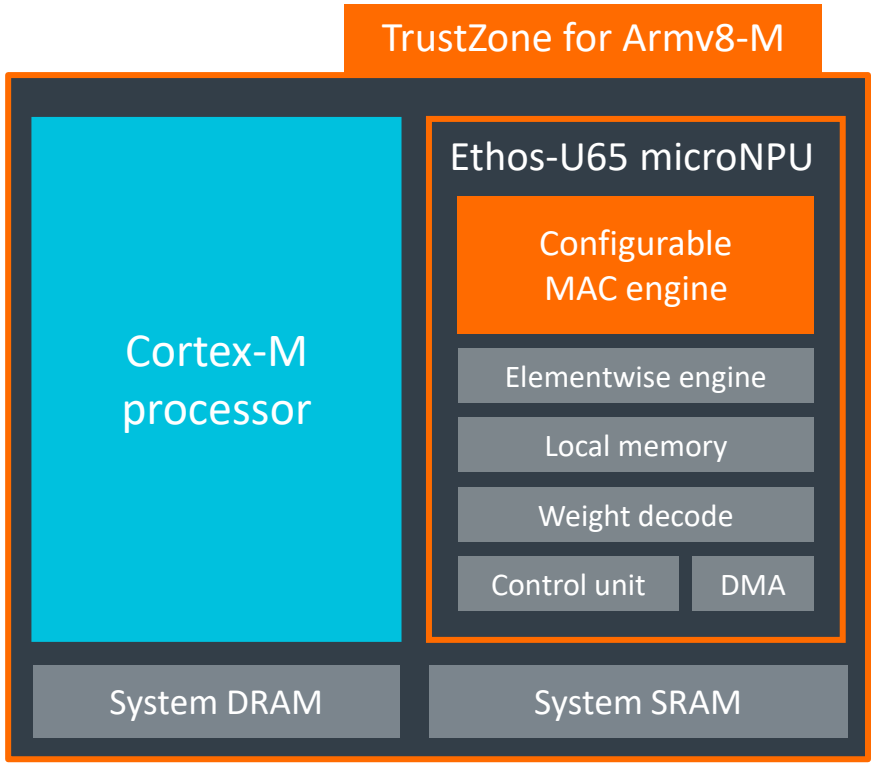

arm

Mapping of NNs to Ethos-U using TensorFlow Lite

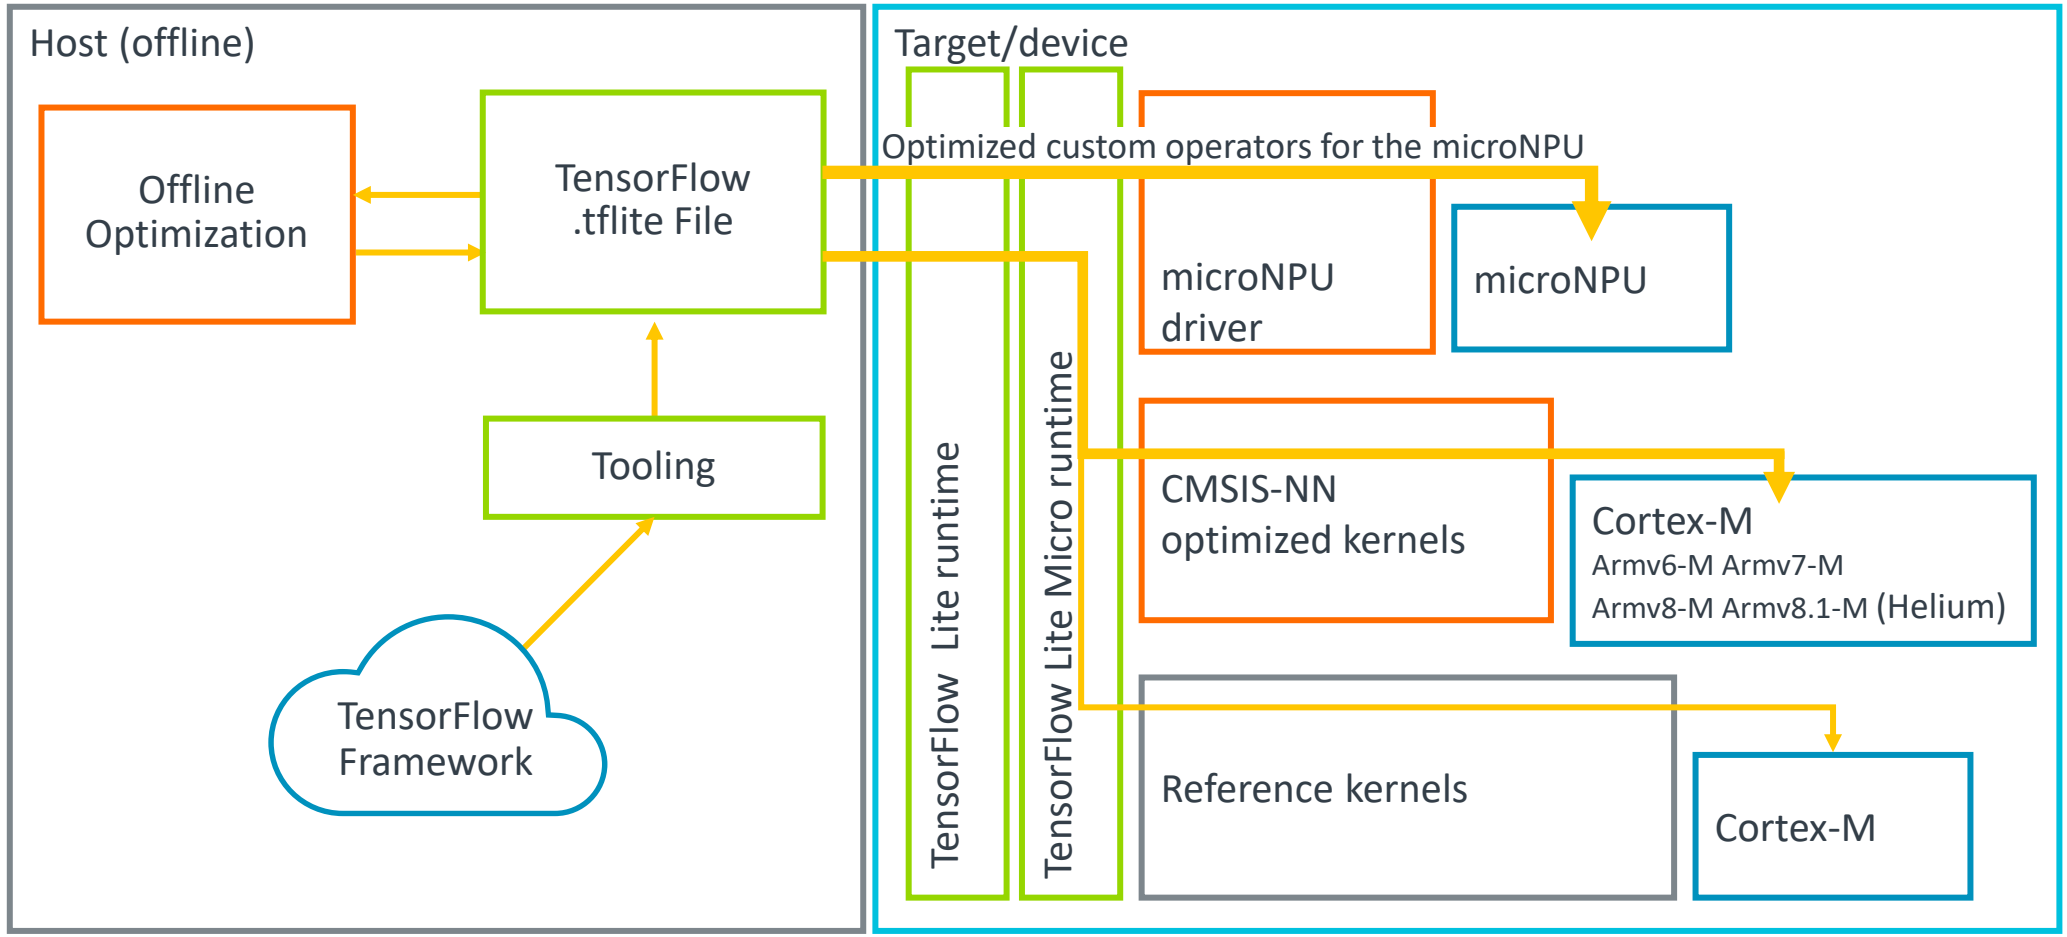

arm

The Vela Optimizer

- Open source compiler
- Reads a tflite file and identifies subgraphs
- Optimizes scheduling of subgraphs
- Loss-less compression of weights
- Generates commands for microNPU
- Writes out a modified tflite file

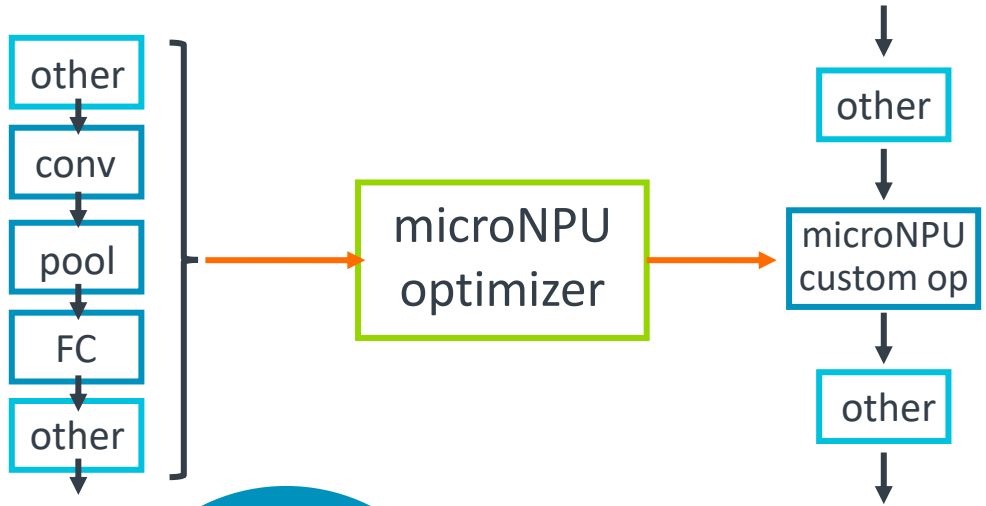

Up to 90% SRAM size reduction

Up to 70% model size reduction

Enabling networks not before feasible in embedded systems

arm

Neural Network performance Across ARM IPs

Wav2Letter

Efficient Software (CMSIS-NN)

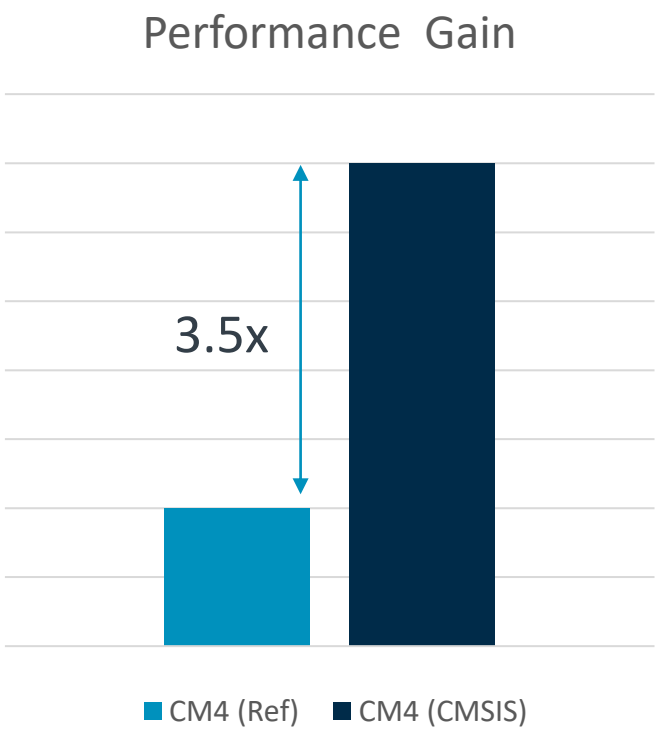

AI Capable Cortex-M55

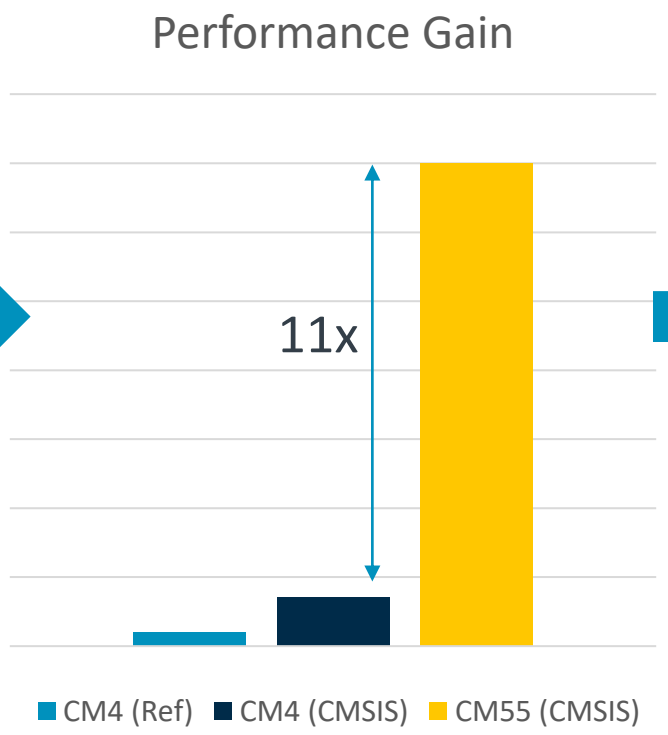

AI Dedicated U55 256 MAC/cycle

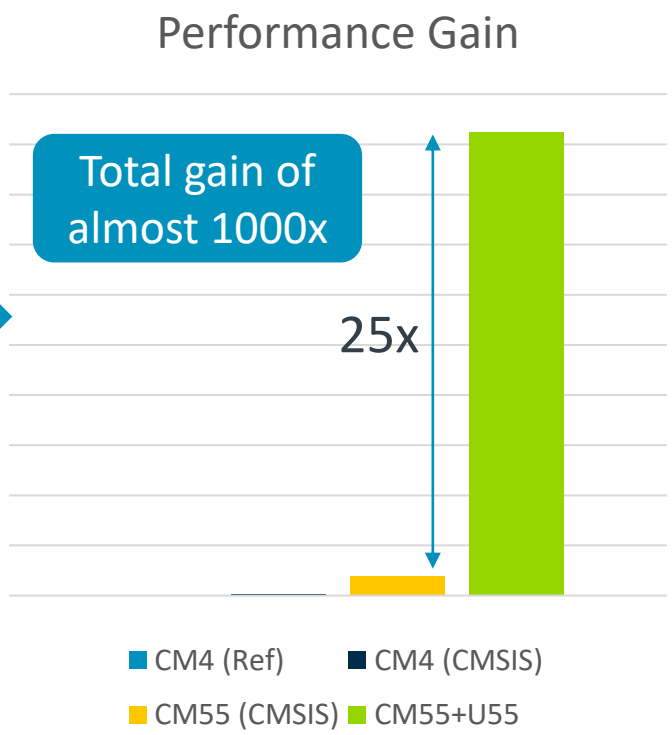

arm

Full example: Typical ML Workload for a Voice Assistant

Speed to inference

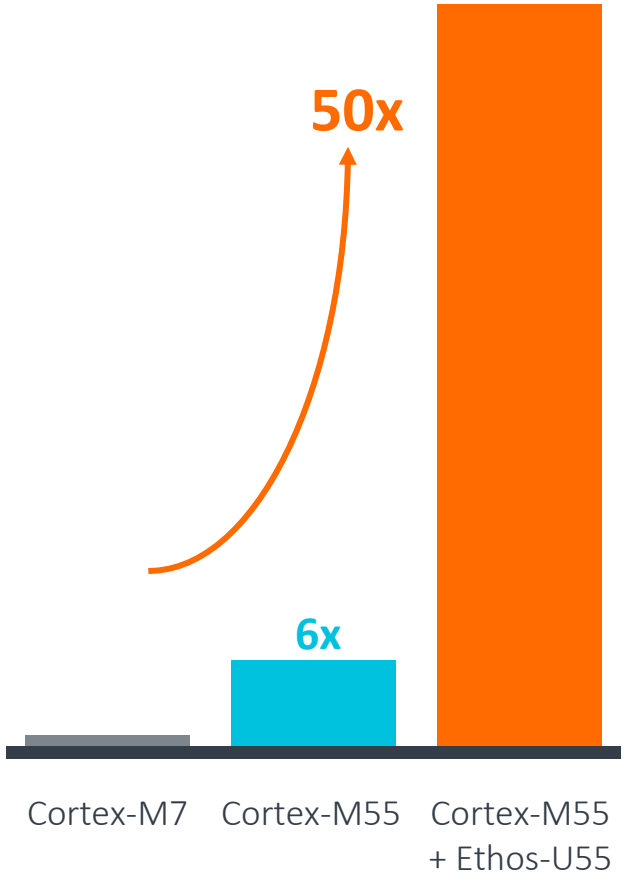

Energy efficiency

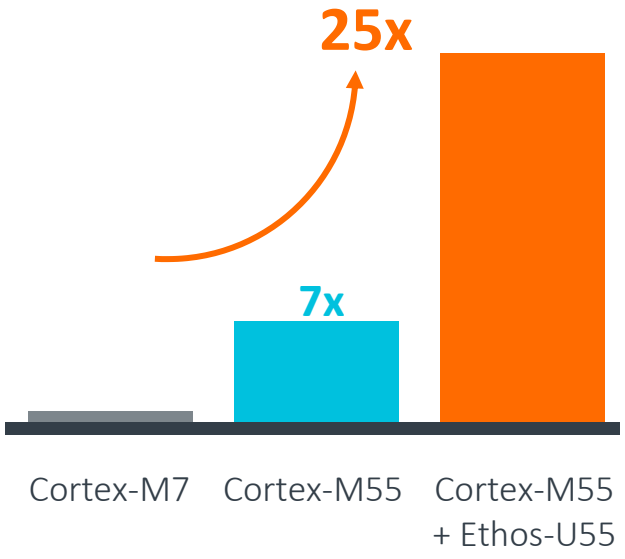

- ✓ Faster responses
- ✓ Smaller form-factors
- ✓ Improved accuracy

Latency and energy spent for all tasks listed combined: voice activity detection, noise cancellation, two-mic beamforming, echo cancellation, equalizing, mixing, keyword spotting, OPUS decode, and automatic speech recognition.

arm

Broadest Range of ML-optimized Processing Solutions

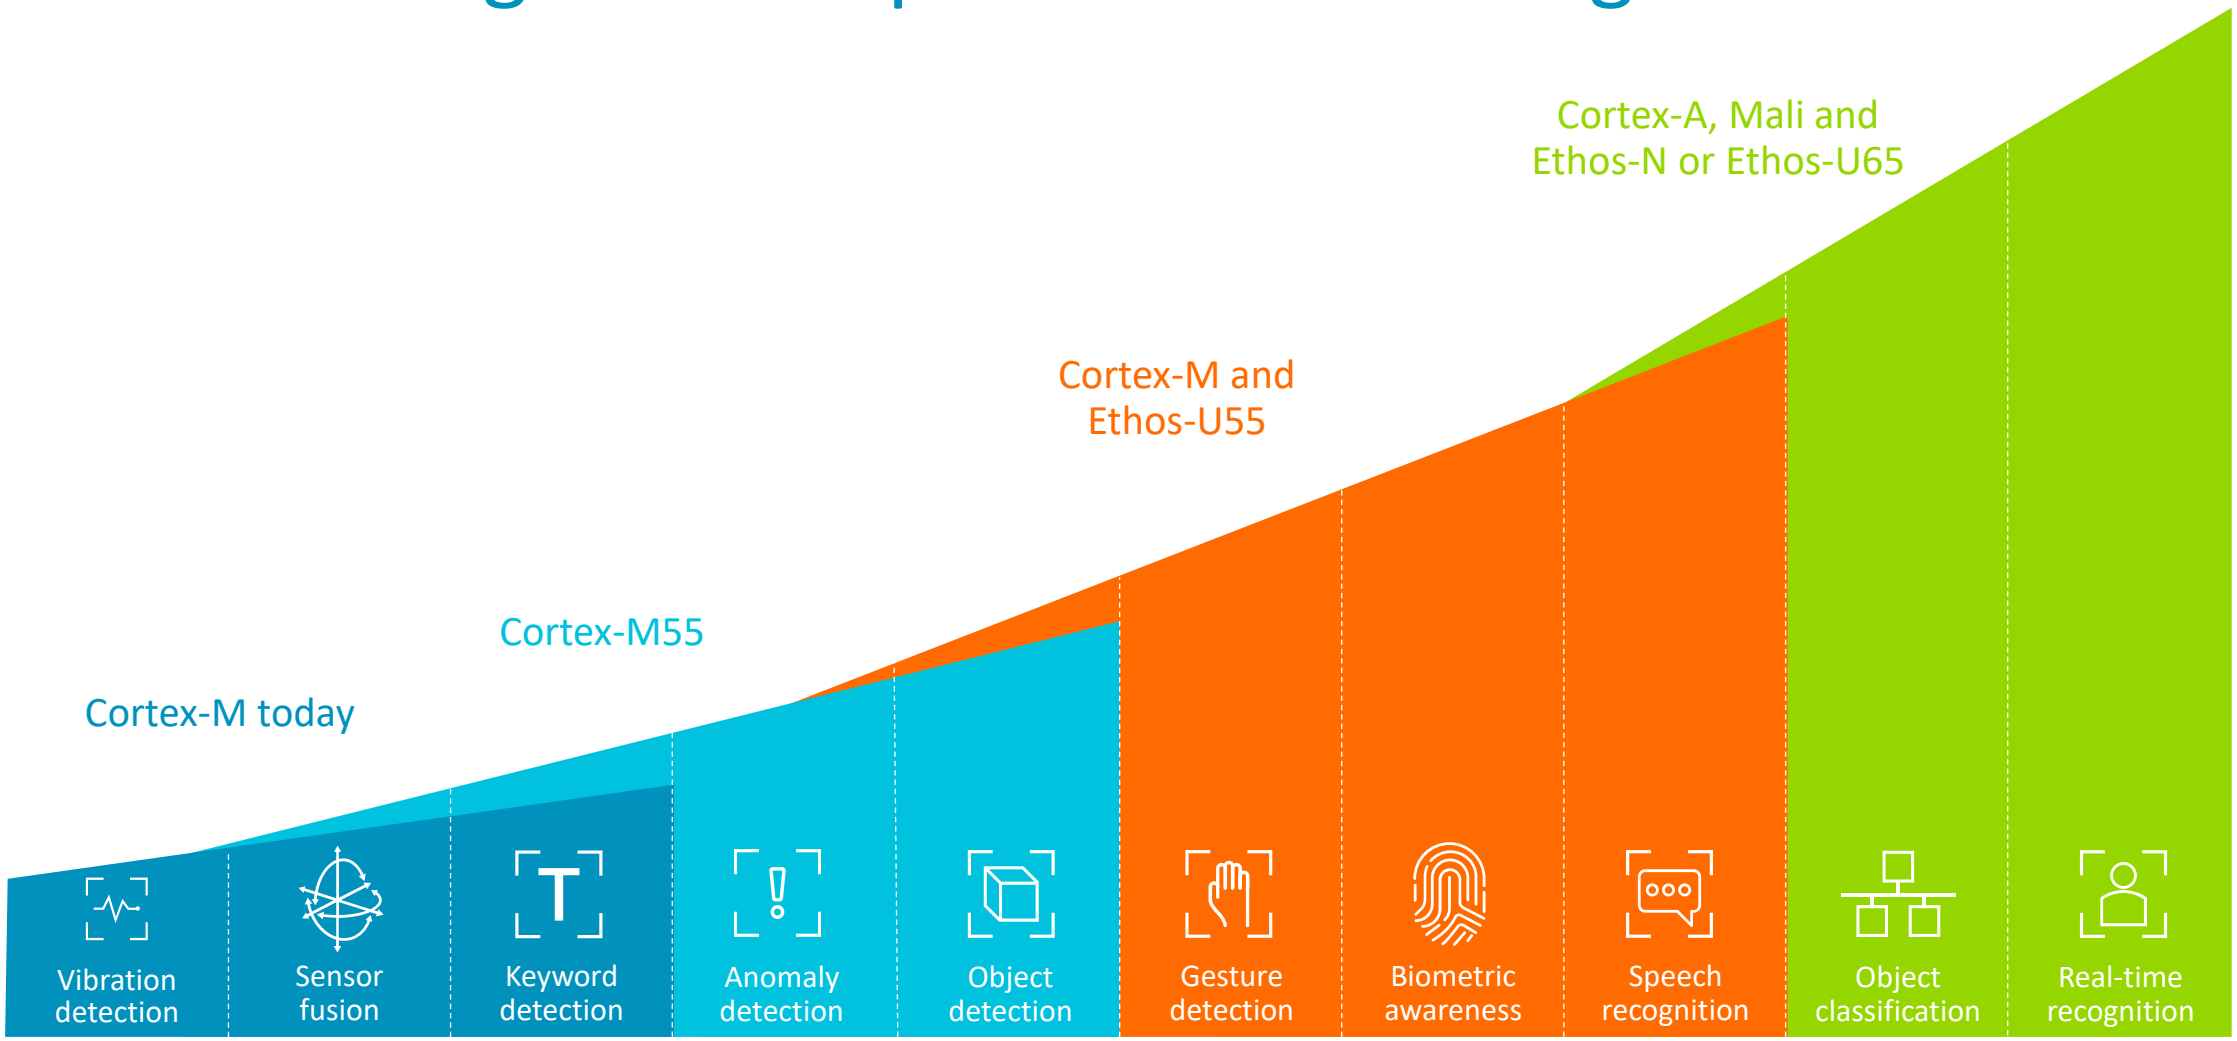

arm